


Compilers without Reinventing a Wheel:

What's up with
MLIR, Mojo 🔥 and beyond

Vasily Ryabov

Expert at Huawei (May 2024)

phd 2X 



~~What is~~ Who is the «Novice»?

Vasily Ryabov

Since 2004 used C++, pure C, Python (2008+), mixed Python/C.

In compilers field –
last 1,5 years, from scratch.



What is a compiler?

Frontend (lexer/tokenizer + syntax)

Source code => tokens (lexemes) => AST (Abstract Syntax Tree)

Parser generators: ANTLR, PEG, ...

Middle end (semantics)

Machine independent IRs (intermediate representations)


Type inference and type checks, common optimizations, exceptions, ...

Backend («90% of the whole job»)

Machine dependent IRs

optimizations => register allocation => Assembler => machine code.

+ JIT (Just-in-Time) engine



What would we do?

Frontend

We will call it from Python 3.9

- in 3.10+ function `PyParser_ASTFromFile(...)` is hidden from C API


Middle end

We will generate IR with help of MLIR framework

We will lower it and pass further

Backend

We will call JIT (Just-in-Time) engine from LLVM



(middle end) What is MLIR?

- Framework «Multi-Level IR» (2019, Chris Lattner & co)
 - High-level and low-level dialects
 - Different dialects in one IR (way to languages interop through IR? ~10+ years)
- LLVM project: [llvm/llvm-project](https://llvm.org/docs/Project.html) (folder mlir/)
- Built on top of MLIR:



Moving to MLIR currently:

clang IR (~30%)

(my sharp estimation: it needs at least 2-3 years)

How such IR looks like? (1/3)

- IR is an SSA form (Static Single Assignment) in both MLIR and LLVM IR
 - each value has unique name/number, it is assigned only once
 - each value can be used many times

Any IR has three representations:


1. In-memory representation ~= C++ classes
2. Bytecode: serialized binary representation in a file (*.bc)
3. Textual IR (*.mlir): not just dump, it could be parsed into in-memory representation

- Functions can't contain in-place functions (like in Python)
- There is nothing similar to `#include` / `import` of other IR module (one IR -> one object file)

How such IR looks like? (2/3)

```
%result = dialect.operation @name (%operands) -> (result_type) { region }
```

```
if True:  
    a = 5  
else:  
    a = 7  
b = a
```



```
func.func @main() -> i64 {  
    %0 = arith.constant 1 : i1  
    %2 = scf.if %0 -> (i64) {  
        %1 = arith.constant 5 : i64  
        scf.yield %1 : i64  
    } else {  
        %1 = arith.constant 7 : i64  
        scf.yield %1 : i64  
    }  
    func.return %2 : i64  
}
```


How such IR looks like? (2/3)

```
%result = dialect.operation @name (%operands) -> (result_type) { region }
```

Some dialects:

- “scf” – structured control flow
- “func” – higher order functions
- “arith” – arithmetic operations

```
if True:
    a = 5
else:
    a = 7
b = a
```




```
func.func @main() -> i64 {
    %0 = arith.constant 1 : i1
    %2 = scf.if %0 -> (i64) {
        %1 = arith.constant 5 : i64
        scf.yield %1 : i64
    } else {
        %1 = arith.constant 7 : i64
        scf.yield %1 : i64
    }
    func.return %2 : i64
}
```


How such IR looks like? (3/3)

<https://godbolt.org/z/j5WeEf8a4>

```
llvm.func @main() -> i64 {  
    %0 = llvm.mlir.constant(true) : i1  
    llvm.cond_br %0, ^bb1, ^bb2  
^bb1:  
    %1 = llvm.mlir.constant(5 : i64) : i64  
    llvm.br ^bb3(%1 : i64)  
^bb2:  
    %2 = llvm.mlir.constant(7 : i64) : i64  
    llvm.br ^bb3(%2 : i64)  
^bb3(%3: i64):  
    llvm.br ^bb4  
^bb4:  
    llvm.return %3 : i64  
}
```



```
func.func @main() -> i64 {  
    %0 = arith.constant 1 : i1  
    %2 = scf.if %0 -> (i64) {  
        %1 = arith.constant 5 : i64  
        scf.yield %1 : i64  
    } else {  
        %1 = arith.constant 7 : i64  
        scf.yield %1 : i64  
    }  
    func.return %2 : i64  
}
```


How different MLIR & LLVM IR?

MLIR диалект "llvm"

```

%0 = llvm.mlir.constant(true) : i1
llvm.cond_br %0, ^bb1, ^bb2
^bb1:
  %1 = llvm.mlir.constant(5 : i64) : i64
  llvm.br ^bb3(%1 : i64)
^bb2:
  %2 = llvm.mlir.constant(7 : i64) : i64
  llvm.br ^bb3(%2 : i64)
^bb3(%3: i64):
  llvm.br ^bb4
^bb4:
  llvm.return %3 : i64

```



LLVM IR

```

br i1 true, label %1, label %2
1:
br label %3
2:
br label %3
3:
%4 = phi i64 [ 7, %2 ], [ 5, %1 ]
br label %5
5:
return %4 : i64

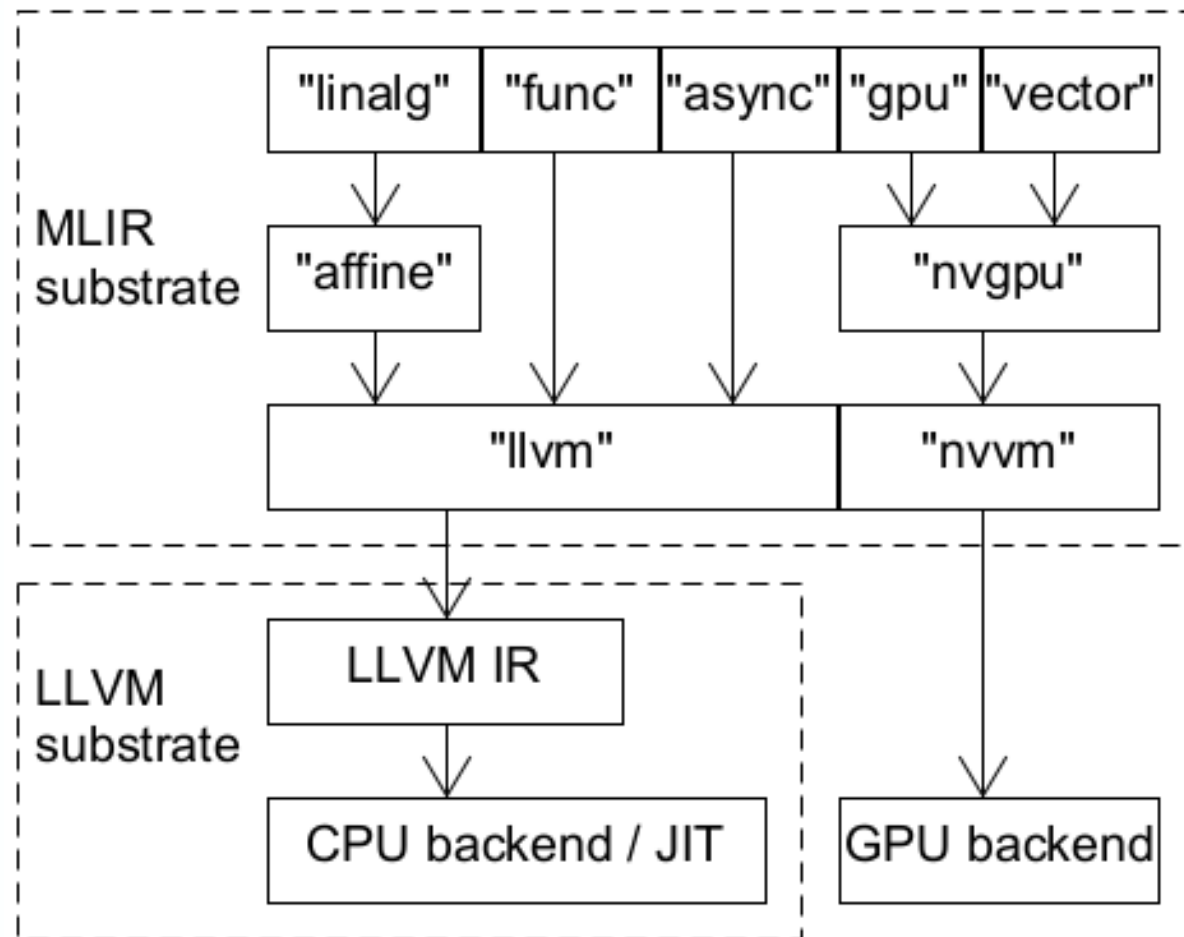
```

How to install MLIR?

- GitHub repository for this talk: [vasily-v-ryabov/phdays24](https://github.com/vasily-v-ryabov/phdays24)
 - Scripts, CMake files and source code
- On Ubuntu/Debian (WSL on Windows) there are apt packages
 - **Already built** LLVM and MLIR libraries and tools 17.x and [18.x](#)
 - [install.sh](#)
- On Windows LLVM repo has to be cloned so that **MLIR is built from sources**
 - MLIR build takes up to 1 hour
 - Visual Studio 2022 must be installed (Community Edition is enough)
 - [install.bat](#) (run in VS2022 command prompt)



What happens inside MLIR?



- [talk about substrates](#)



- 1. Generating high-level IR**
(by visiting/walking Python AST)
- 2. Lowering step by step**
(Lowering/Conversion is a part of dialect)
- 3. Translation to LLVM IR**
- 4. Execution by LLVM's backend**

How will we generate IR?

“Toy” project: [llvm/llvm-project: ./mlir/examples/toy/](https://llvm.org/docs/GettingStarted.html#_toy_project) (7 chapters)

- 1) frontend, **2) MLIR gen**, **3) dialect**, 4) pass, 5) own lowering pass, **6) ->LLVM IR**, **7) JIT**.

[vasily-v-ryabov/phdays24](https://github.com/vasily-v-ryabov/phdays24) has 5 chapters:

- | | |
|--|-------------------|
| 1) The simplest IR generation | (22 lines in C++) |
| 2) IR for variables assignments (with Python 3.9 frontend) | (+199 lines) |
| 3) IR with conditions and type inference | (+59 lines) |
| 4) Lowering and translation to LLVM IR (with options) | (+49 lines) |
| 5) Using JIT engine | (+34 lines) |
| | Total: 363 lines |

How will we generate IR?

vasily-v-ryabov/phdays24 has 5 chapters:

- 1) The simplest IR generation**
- 2) IR for variables assignments (with Python 3.9 frontend)
- 3) IR with conditions and type inference
- 4) Lowering and translation to LLVM IR (with options)
- 5) Using JIT engine

Ya! The first code: main.cpp (1/3)



```
#include "mlir/IR/MLIRContext.h"
```

```
#include "mlir/IR/Verifier.h"
```

[./phdays24/01_MLIR_gen/main.cpp](https://github.com/llvm/llvm-project/blob/main/mlir/examples/phd2x/main.cpp)

```
#include "mlir/Dialect/LLVMIR/LLVMDialect.h"
```

```
mlir::ModuleOp mlirGen(mlir::MLIRContext &context) {  
    mlir::OpBuilder builder(&context);  
    context.getOrLoadDialect<mlir::LLVM::LLVMDialect>();  
    ...  
    auto loc = builder.getUnknownLoc();  
  
    auto module = mlir::ModuleOp::create(loc);  
    builder.setInsertionPointToEnd(module.getBody());  
    ...  
}
```

Ya! The first code: main.cpp (2/3)



```
...  
// create function main()  
auto mainFuncType = mlir::LLVM::LLVMFunctionType::get(builder.getI32Type(), {});  
auto mainFunc = builder.create<mlir::LLVM::LLVMFuncOp>(loc, "main", mainFuncType);  
  
mlir::Block *entryBlock = mainFunc.addEntryBlock();  
builder.setInsertionPointToStart(entryBlock);  
  
// function body (return 0;)  
auto constOp = builder.create<mlir::LLVM::ConstantOp>(loc, builder.getI32Type(), 0);  
builder.create<mlir::LLVM::ReturnOp>(loc, constOp->getResult(0));  
return module;  
}
```


Ya! The first code: main.cpp (3/3)



```
int main(int argc, char **argv) {
    mlir::MLIRContext context;
    mlir::OwningOpRef<mlir::ModuleOp> module = mlirGen(context);
    if (!module)
        return 1;
    if (mlir::failed(mlir::verify(*module))) {
        module->emitError("Module verification failed!");
        module->dump(); // dump incorrect IR anyway
        return 2;
    }
    module->dump(); // to stderr
    return 0;
} // 22 non-empty lines of code
```

Ya! The first code: main.cpp (3/3)



```
int main(int argc, char **argv) {
    mlir::MLIRContext context;
    mlir::OwningOpRef<mlir::ModuleOp> module = mlirGen(context);
    if (!module)
        return 1;
    if (mlir::failed(mlir::verify(*module))) {
        module->emitError("Module verification failed!");
        module->dump(); // dump incorrect IR anyway
        return 2;
    }
    module->dump(); // to stderr
    return 0;
} // 22 non-empty lines of code
```

```
root@MSI:~/phdays24/01_MLIR_gen/build# ./bin/py39compiler
module {
  llvm.func @main() -> i32 {
    %0 = llvm.mlir.constant(0 : i32) : i32
    llvm.return %0 : i32
  }
}
```

How will we generate IR?



```
a = 5
b = 7
b = a
```

vasily-v-ryabov/phdays24 has 5 chapters:

- 1) The simplest IR generation
- 2) IR for variables assignment** (with Python 3.9 frontend)
- 3) IR with conditions and type inference
- 4) Lowering and translation to LLVM IR (with options)
- 5) Using JIT engine



How to support variables? (1/4)



[./phdays24/02_MLIR_gen_pyvars/include/py_ast.h](https://phdays24/02_MLIR_gen_pyvars/include/py_ast.h)

- class-wrapper for function RýRắssês AŞŦGsộñGîlê

```
a = 5
b = 7
b = a
```

```
class PyAST {
...
    bool parse_file(const char *name) { ... }
    mod_ty mod() { ... } // mod_ty is a pointer to C structure in Python.h
};
```

How to support variables? (2/4)

[./phdays24/02_MLIR_gen_pyvars/include/MLIRGen.h](#)

```
mlir::LogicalResult mlirGen(mod_ty pyModule) { ... }
mlir::LogicalResult mlirGen(asdl_seq *statements) { ... }
mlir::LogicalResult mlirGen(stmt_ty statement) {
    switch (statement->kind) {
        ...
        case Assign_kind: {
            auto valOrErr = mlirGen(/*expr_ty*/statement->v.Assign.value); // right side
            ... // left side
            expr_ty astTarget = (expr_ty)asdl_seq_GET(statement->v.Assign.targets, 0);
            ...
        }
    }
    mlir::FailureOr<mlir::Value> mlirGen(expr_ty expr) { ... } // case Constant_kind:
```

How to support variables? (3/4)

```
a = 5
```

```
b = 7
```

```
b = a
```

- At the right side there is **name**, but where to get **mlir::Value** ?
- We need a symbol table! But it's not a simple `std::hash_map`...

```
#include "llvm/ADT/ScopedHashTable.h"
```

```
...
```

```
llvm::ScopedHashTable<llvm::StringRef, mlir::Value> symbolTable;
```

```
...
```

```
// at least one scope is required
```

```
llvm::ScopedHashTableScope<llvm::StringRef, mlir::Value> scope(symbolTable);
```

How to support variables? (4/4)

```
void defineVariable(llvm::StringRef name, mlir::Value value) {  
    llvm::outs() << "Add variable '" << name << "' = '" << value << "'\n";  
    llvm::MallocAllocator ma;  
    symbolTable.insert(name.copy(ma), value);  
}  
  
mlir::FailureOr<mlir::Value> getVariable(mlir::Location loc, llvm::StringRef name) {  
    auto value = symbolTable.lookup(name);  
    if (value) // may be nullptr!  
        return value;  
    mlir::emitError(loc, "Variable '" << name << "' is not defined\n");  
    return mlir::failure();  
}
```

- **MLIRGen.h**: ~200 lines of code, folder in the repo: **02_MLIR_gen_pyvars/**
- **py_ast.h**: ~50 lines of code

How to support variables? (4/4)

```
void defineVariable(llvm::StringRef name, mlir::Value value) {
    llvm::outs() << "Add variable '" << name << "' = '" << value << "'\n";
    llvm::MallocAllocator ma;
    symbolTable.insert(name.copy(ma), value);
}
```

```
root@MSI:~/phdays24/02_MLIR_gen_pyvars# ./build/bin/py39compiler script.py
Add variable 'a' = '%0 = "llvm.mlir.constant"() <{value = 5 : i64}> : () -> i64'
Add variable 'b' = '%1 = "llvm.mlir.constant"() <{value = 7 : i64}> : () -> i64'
Add variable 'b' = '%0 = "llvm.mlir.constant"() <{value = 5 : i64}> : () -> i64'
module {
  llvm.func @main() -> i32 {
    %0 = llvm.mlir.constant(5 : i64) : i64
    %1 = llvm.mlir.constant(7 : i64) : i64
    %2 = llvm.mlir.constant(0 : i32) : i32
    llvm.return %2 : i32
  }
}
```


What about globals?

- 1) Create LLVM::GlobalOp ("llvm.mlir.global") – it is always a pointer
 - 2) Get address by name: LLVM::AddressOfOp ("llvm.addressof")
 - 3) Load from pointer: LLVM::LoadOp ("llvm.load")
 - 4) Or store to pointer: LLVM::StoreOp ("llvm.store")
- No scopes at all!
 - "llvm.mlir.global" is similar to insert, but at runtime
 - "llvm.addressof"+"llvm.load" are similar to lookup, but at runtime
 - Functions are also globals
 - There is `mlir::SymbolTable` `globalTable(module)`;
 - There is insert (not into IR) and lookup (useful to generate function body later)
 - There is `module.lookupSymbol(name)`;
 - If "llvm.addressof" returned a function, it can be called by pointer (indirect call)

How will we generate IR?



```
if True:
    a = 5
else:
    a = 2
b = a
```

vasily-v-ryabov/phdays24 has 5 chapters:

- 1) The simplest IR generation
- 2) IR for variables assignment (with Python 3.9 frontend)
- 3) IR with conditions and type inference**
- 4) Lowering and translation to LLVM IR (with options)
- 5) Using JIT engine



How will we add conditions? (1/5) ~~phd 2~~

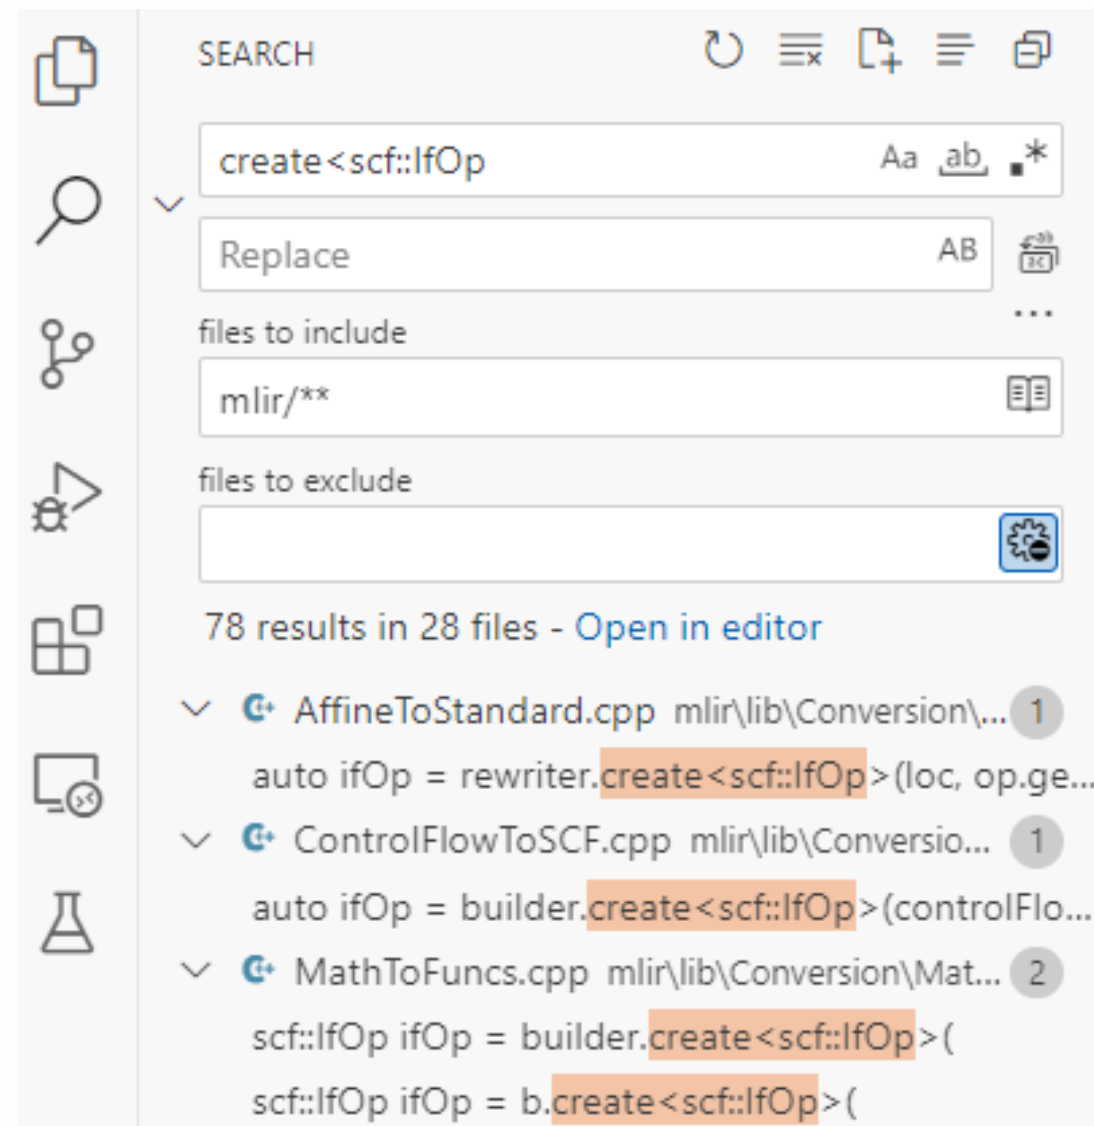
```
#include "mlir/Dialect/SCF/IR/SCF.h"
#include "mlir/Dialect/ControlFlow/IR/ControlFlow.h"
...
context.getOrLoadDialect<mlir::scf::SCFDialect>();
context.getOrLoadDialect<mlir::cf::ControlFlowDialect>();
...
```

```
auto ifOp = builder.create<mlir::scf::IfOp>(loc, ...); // parameters?
```

```
if True:
    a = 5
else:
    a = 2
b = a
```

How to find builders

- No builders' info in the dialect docs!
- We have to find them in the MLIR sources



How to find builders

- No builders' info in the dialect docs!
- We have to find them in the MLIR sources
- If we go into class `scf::IfOp`, ...

```
static void build(::mlir::OpBuilder &odsBuilder, ::mlir::OperationState &odsState, TypeRange resultTypes, Value cond);  
static void build(::mlir::OpBuilder &odsBuilder, ::mlir::OperationState &odsState, TypeRange resultTypes, Value cond, bool addThenBlock, bool addElseBlock);  
static void build(::mlir::OpBuilder &odsBuilder, ::mlir::OperationState &odsState, Value cond, bool withElseRegion);  
static void build(::mlir::OpBuilder &odsBuilder, ::mlir::OperationState &odsState, TypeRange resultTypes, Value cond, bool withElseRegion);  
static void build(::mlir::OpBuilder &odsBuilder, ::mlir::OperationState &odsState, Value cond, function_ref<void(OpBuilder &, Location)> thenBuilder = buildT
```

How to find builders

- No builders' info in the dialect docs!
- We have to find them in the MLIR sources
- If we go into class `scf::IfOp`, ...

```
, TypeRange resultTypes, Value cond);  
, TypeRange resultTypes, Value cond, bool addThenBlock, bool addElseBlock);  
, Value cond, bool withElseRegion);  
, TypeRange resultTypes, Value cond, bool withElseRegion);  
, Value cond, function_ref<void(OpBuilder &, Location)> thenBuilder = buildTerminatedBody,
```

How will we add conditions? (2/5)

```

auto valueOrError = mlirGen(statement->v.If.test);
if (mlir::failed(valueOrError))
    return mlir::failure();
auto condition = valueOrError.value(); // TODO: check type is i1 (integer of 1 bit)
auto ifOp = builder.create<mlir::scf::IfOp>( // it calls inferReturnTypes(...)
    loc, condition,
    /*thenBuilder=*/[&](mlir::OpBuilder &b, mlir::Location loc) {

        b.create<mlir::scf::YieldOp>(loc /*returned values*/);
    },
    /*elseBuilder=*/[&](mlir::OpBuilder& b, mlir::Location loc) {

        b.create<mlir::scf::YieldOp>(loc /*returned values*/);
    }
    );

```

How will we add conditions? (2/5)

```
auto valueOrError = mlirGen(statement->v.If.test);
if (mlir::failed(valueOrError))
    return mlir::failure();
auto condition = valueOrError.value(); // TODO: check type is i1 (integer of 1 bit)
auto ifOp = builder.create<mlir::scf::IfOp>( // it calls inferReturnTypes(...)
    loc, condition,
    /*thenBuilder=*/[&](mlir::OpBuilder &b, mlir::Location loc) {
        mlirGen(statement->v.If.body); // Assign_kind AST-nodes can be inside
        b.create<mlir::scf::YieldOp>(loc /*returned values*/);
    },
    /*elseBuilder=*/[&](mlir::OpBuilder& b, mlir::Location loc) {
        mlirGen(statement->v.If.orelse); // Assign_kind AST-nodes can be inside
        b.create<mlir::scf::YieldOp>(loc /*returned values*/);
    }
);
```


How will we add conditions? (3/5) ~~Phd 2~~

```
std::stack<std::set<llvm::StringRef>> ifElseVariables;
llvm::MallocAllocator ma; // for copying llvm::StringRef's
...
void defineVariable(llvm::StringRef name, mlir::Value value) {
    symbolTable.insert(name.copy(ma), value);
    if (!ifElseVariables.empty())
        ifElseVariables.top().insert(name.copy(ma));
}
```

How will we add conditions? (4/5) ~~Phd 2~~

```
/*thenBuilder=*/  
[&](mlir::OpBuilder &b, mlir::Location loc) {  
    ifElseVariables.push(std::set<llvm::StringRef>());  
    result = mlirGen(statement->v.If.body);  
    auto varsSet = ifElseVariables.top();  
    llvm::SmallVector<mlir::Value> returnValues;  
    for (auto it = varsSet.begin(); it != varsSet.end(); it++) {  
        auto value = symbolTable.lookup(*it);  
        if (value)  
            returnValues.push_back(value);  
    }  
    ifElseVariables.pop();  
    b.create<mlir::scf::YieldOp>(loc, returnValues);  
},
```

How will we add conditions? (5/5) ~~phd 2~~

```
if True:
```

```
    a = 5
```

```
else:
```

```
    a = 2
```

```
b = a
```

(slide 7 → slide 35) == 🔥

```
root@MSI:~/phdays24/03_MLIR_gen_if_else# ./build/bin/py39compiler script.py
module {
  llvm.func @main() -> i32 {
    %0 = llvm.mlir.constant(true) : i1
    %1 = scf.if %0 -> (i64) {
      %3 = llvm.mlir.constant(5 : i64) : i64
      scf.yield %3 : i64
    } else {
      %3 = llvm.mlir.constant(2 : i64) : i64
      scf.yield %3 : i64
    }
    %2 = llvm.mlir.constant(0 : i32) : i32
    llvm.return %2 : i32
  }
}
```

How will we generate IR?

vasily-v-ryabov/phdays24 has 5 chapters:

- 1) The simplest IR generation
- 2) IR for variables assignment (with Python 3.9 frontend)
- 3) IR with conditions and type inference
- 4) Lowering and translation to LLVM IR (with options)**
- 5) Using JIT engine



How to run lowering

1. "scf" – structured control flow • `mlir-opt-18 --help | grep to-llvm`
2. "cf" – control flow • `mlir-opt-18 --convert-scf-to-cf input.mlir`
3. "llvm" • `mlir-opt-18 --convert-cf-to-llvm input.mlir`

How to run lowering

1. "scf" – structured control flow • `mlir-opt-18 --help | grep to-llvm`
2. "cf" – control flow • `mlir-opt-18 --convert-scf-to-cf input.mlir`
3. "llvm" • `mlir-opt-18 --convert-cf-to-llvm input.mlir`
4. LLVM IR • `mlir-translate-18 --mlir-to-llvmir input.mlir`

How to run lowering

1. "scf" – structured control flow
 2. "cf" – control flow
 3. "llvm"
 4. LLVM IR
 5. Execute on JIT engine
- `mlir-opt-18 --help | grep to-llvm`
 - `mlir-opt-18 --convert-scf-to-cf input.mlir`
 - `mlir-opt-18 --convert-cf-to-llvm input.mlir`
 - `mlir-translate-18 --mlir-to-llvmir input.mlir`
 - `lli output.ll` or `mlir-cpu-runner input.mlir`

How to run lowering

1. "scf" – structured control flow
 2. "cf" – control flow
 3. "llvm"
 4. LLVM IR
 5. Execute on JIT engine
- `mlir-opt-18 --help | grep to-llvm`
 - `mlir-opt-18 --convert-scf-to-cf input.mlir`
 - `mlir-opt-18 --convert-cf-to-llvm input.mlir`
 - `mlir-translate-18 --mlir-to-llvmir input.mlir`
 - `lli output.ll` or `mlir-cpu-runner input.mlir`
- (what if it crashed? w/ huge IR!)
- `mlir-opt-18 --mlir-pass-pipeline-crash-reproducer=<output_filepath.mlir>`

How to run lowering

1. "scf" – structured control flow
 2. "cf" – control flow
 3. "llvm"
 4. LLVM IR
 5. Execute on JIT engine
- `mlir-opt-18 --help | grep to-llvm`
 - `mlir-opt-18 --convert-scf-to-cf input.mlir`
 - `mlir-opt-18 --convert-cf-to-llvm input.mlir`
 - `mlir-translate-18 --mlir-to-llvmir input.mlir`
 - `lli output.ll` or `mlir-cpu-runner input.mlir`

(what if it crashed? w/ huge IR!)

- `mlir-opt-18 --mlir-pass-pipeline-crash-reproducer=<output_filepath.mlir>`

(how to get back to MLIR?)

- 4. LLVM IR → 3. MLIR "llvm"
- `clang -S -emit-llvm foo.c >foo.ll`
- `mlir-translate-18 --import-llvm foo.ll >foo.mlir`

How to write lowering

```
#include "mlir/Conversion/SCFToControlFlow/SCFToControlFlow.h"
#include "mlir/Conversion/ControlFlowToLLVM/ControlFlowToLLVM.h"
#include "mlir/Pass/Pass.h"
#include "mlir/Pass/PassManager.h"

...

int main(int argc, char **argv) {
    ...
    mlir::PassManager passes(&context);
    passes.addPass(mlir::createConvertSCFToCFPass());
    passes.addPass(mlir::createConvertControlFlowToLLVMPass());
    if (mlir::failed(passes.run(module.get())))
        return 5;
}
```

How to translate to LLVM IR

```
#include "mlir/Target/LLVMIR/Dialect/Builtin/BuiltinToLLVMIRTranslation.h"
#include "mlir/Target/LLVMIR/Dialect/LLVMIR/LLVMToLLVMIRTranslation.h"
#include "mlir/Target/LLVMIR/Export.h"

...

int main(int argc, char **argv) {
    ...
    mlir::registerBuiltinDialectTranslation(*module->getContext());
    mlir::registerLLVMDialectTranslation(*module->getContext());

    llvm::LLVMContext llvmContext;
    auto llvmModule = mlir::translateModuleToLLVMIR(*module, llvmContext);
    llvm::errs() << *llvmModule << "\n"; // dump LLVM IR
}
```

How to parse CmdLine options

```
#include "llvm/Support/CommandLine.h"

...

enum Action { DumpMLIR, DumpLLVM, DumpLLVMIR };
namespace cl = llvm::cl;

static cl::opt<enum Action> emitAction("emit", cl::desc("Select the output"),
    cl::values(clEnumValN(DumpMLIR, "mlir", "output the MLIR dump")),
    cl::values(clEnumValN(DumpLLVM, "llvm", "dump the MLIR \"llvm\" dialect")),
    cl::values(clEnumValN(DumpLLVMIR, "llvm-ir", "output the LLVM IR dump")));

static cl::opt<std::string> srcPy(cl::Positional, cl::desc("<input .py file>"),
    cl::init("-"), cl::value_desc("filename"));

int main(int argc, char **argv) {
    cl::ParseCommandLineOptions(argc, argv, "Python 3.9 demo compiler\n");
}
```

How to parse CmdLine options

```
root@MSI:~/phdays24/04_MLIR_gen_LLVM_IR# ./build/bin/py39compiler -emit=llvm-ir script.py
; ModuleID = 'LLVMDialectModule'
source_filename = "LLVMDialectModule"

define i32 @main() {
    br i1 true, label %1, label %2

1:                                     ; preds = %0
    br label %3

2:                                     ; preds = %0
    br label %3

3:                                     ; preds = %1, %2
    %4 = phi i64 [ 2, %2 ], [ 5, %1 ]
    br label %5

5:                                     ; preds = %3
    ret i32 0
}

!llvm.module.flags = !{!0}

!0 = !{i32 2, !"Debug Info Version", i32 3}
```

How will we generate IR?

vasily-v-ryabov/phdays24 has 5 chapters:

- 1) The simplest IR generation
- 2) IR for variables assignment (with Python 3.9 frontend)
- 3) IR with conditions and type inference
- 4) Lowering and translation to LLVM IR (with options)
- 5) **Using JIT engine**



How to integrate JIT engine

```
#include "mlir/ExecutionEngine/ExecutionEngine.h"
#include "mlir/ExecutionEngine/OptUtils.h"
#include "llvm/ExecutionEngine/Orc/JITTargetMachineBuilder.h"
#include "llvm/Support/TargetSelect.h"

llvm::InitializeNativeTarget();
llvm::InitializeNativeTargetAsmPrinter();
mlir::ExecutionEngineOptions engineOptions;
auto maybeEngine = mlir::ExecutionEngine::create(*module, engineOptions);
auto &engine = maybeEngine.get();
...
llvm::SmallVector<void *> argsAndReturn; // int32_t main() without argc, argv[]
int32_t exitCode; argsAndReturn.push_back(&exitCode); // address of return value
auto invocationResult = engine->invokePacked("main", argsAndReturn);
```

What else has the JIT engine?

- Hot code detection and its re-optimization
 - example: -O0 is the default level, -O2 is only for hot code
- JIT callback: runtime can call a compiler back (runtime -> compiler -> runtime)
 - example is somewhere in ./llvm/unittests/ExecutionEngine/Orc/
- Support of the debugger for your language (ORCDebugging library)
- Materialization layers (used by Mojo 🔥 a lot)
- <https://www.llvm.org/docs/ORCv2.html>

What is absent in “scf” dialect?

- “scf.break”
- “scf.continue”
- early return
 - “func.return”: HasParent<FuncOp> (terminator only for “func.func”)
 - “llvm.return”: it will be lowered as is, but we need “llvm.br” to a final block
- But we can:
 - use dialect “cf”: “cf.br” и “cf.cond_br” (or “llvm.br” and “llvm.cond_br”)

What's up with Mojo 🔥 ?

Mojo 🔥 has dialect "hlcf": [\(YouTube\) 2023 LLVM Dev mtg Mojo \(time=17:00\)](#)
(High-Level Control Flow)

- "hlcf.**break**"

```
hlcf.loop {
  %1 = lit.ref.load %i : <mut !Int, *"`i0">
  %2 = kgen.param.constant: !Int = <{value = 10}>
  %3 = kgen.call @Int::@__lt__(%1, %2)
  %4 = kgen.call @Bool::@__mlir_i1__(%3)
  hlcf.if %4 {
    hlcf.yield
  } else {
    hlcf.break
  }
}
```

- "hlcf.**continue**"

```
hlcf.loop (%arg2 = %idx0 : index) {
  %0 = index.cmp slt(%arg2, %idx10)
  hlcf.if %0 {
    hlcf.yield
  } else {
    hlcf.break
  }

  %1 = kgen.call @print(%arg2)
  %2 = index.add %arg2, %idx1
  hlcf.continue %2 : index
}
```

- early **return**: it is supported, but without many details

What's up with clangIR?



<clang/include/clang/CIR/Dialect/IR/CIROps.td#L797>

- "cir.**break**"
- "cir.**continue**"
- "cir.**return**" (early return): <clang/include/clang/CIR/Dialect/IR/CIROps.td#L540>
 - ParentOneOf<["FuncOp", "ScopeOp", "IfOp", "SwitchOp", "DoWhileOp", "WhileOp", "ForOp"]



What is *absent* in “func” dialect?



- “func.invoke”
 - it is a function call, which may raise an exception
 - there is “llvm.invoke” (in theory the “llvm” dialect should be enough for exception handling, but a lot of code is required)
- But “func.func” can return multiple values (“llvm.func” can return only one)
 - the workaround is to return “llvm.struct” type (example: “llvm.struct “A”<i64,i64,f64>”)
- If we look into clangIR operations (CIROps.td):
 - “cir.try_call” (cir::TryCallOp), “cir.try” (cir::TryOp)
 - “cir.catch” (cir::CatchOp), “cir.catch_param” (cir::CatchParamOp)
 - “cir.alloc_exception” (cir::AllocExceptionOp), “cir.throw” (cir::ThrowOp)
 - “cir.stack_save” (cir::StackSaveOp), “cir.stack_restore” (cir::StackRestoreOp)
 - (CIRTypes.td) see type “cir.eh.info”
- Class support is required for exceptions (dialect “cir” supports classes)

Which dialects are used by Mojo 🔥 ?

- Upstream dialects "llvm" и "index"
- "hlcf" is High level control flow
- "mosh" is "Mojo 🔥 Shape" dialect (shape is a size of vector, matrix or tensor)
- "kgen" is a meta dialect, it has **ElaborationPass** which do the most of job!
 - for functions with [type params] like this: `var a = func[t1, t2](arg1, arg2)`
 - for structures with [type params] similar to C++ templates
 - for compile time introspection

What is unique in Mojo ?

- **AST** (Abstract Syntax Tree) **is not built** at all!
 - Source code parser immediately generates an MLIR IR
 - MLIR has constant folding hooks by design (typical AST level optimization)
 - [Interview with Chris Lattner \(May 2024\)](#)
- Standard module TargetInfo (low level) is used for loop unrolling at high level
- JIT engine ORC (from LLVM) is used **at all stages** for adaptive compilation
 - both just-in-time (REPL), **and ahead-of-time!**
 - materialization layers are created

What Mojo took from LLVM?

[\(c\) Talk about Mojo in 2023 \(at 27:15\)](#)

80th level of parallelism (at 27:40):

- One LLVMContext per function

LLVM is good for:

- GVN, Load/Store Optimization, LSR, etc
- scalar optimization (e.g. instcombine)
- target-specific code generation

We need to disable:

- Vectorizer, loop unroller, etc
- Inliner and other IPO passes

Solution: replace these!

- Build new MLIR passes
- Replace others with Mojo libraries

Why not learn Mojo yet?

The language is changing regularly:


- let variables were removed recently
- dynamic types at runtime are on re-design

The MLIR dialects are not open:

- "kgen", "hlcf", ... (not mature?)

Priorities (in my opinion at the moment):

- P1: AI engine
- P2: MLIR frontend (the language for compiler engineers)
- P3: Python superset (superset++, i.e. much wider than Python)



Conclusion: learn MLIR and JIT engine ORC => fast learn Mojo 

What else to watch?



← [playlist](#)

In Russian:

- [Constantine Vladimirov's lecture about LLVM IR \(2019\)](#) – GEP, load, alloca

In English:

- [Extending Dominance to MLIR Regions \(2023\)](#)
- [MLIR Dialect Design \(EuroLLVM 2023\)](#) – dialects classification
- [What's new in MLIR? \(2023 vs 2019\)](#)
- [JITLink: Native Windows JITing in LLVM](#) & [ORCv2 - LLVM JIT APIs Deep Dive](#)
- www.youtube.com/@LLVMPROJ – LLVM Dev meetings & EuroLLVM



@vasily_v_ryabov



vasily.v.ryabov@gmail.com



github.com/vasily-v-ryabov



phd 2 Positive Hack Days Fest
от positive technologies



Thank you!



OK build MLIR (backup slide)

```
git clone --depth 1 -b llvmorg-18.1.5 https://github.com/llvm/llvm-project.git
cd ./llvm-project/ && mkdir build && cd build/
pip install -U cmake ninja
cmake -G Ninja ../llvm/ -DLLVM_ENABLE_PROJECTS="mlir" -
DLLVM_ENABLE_ASSERTIONS=ON -DCMAKE_BUILD_TYPE=Debug -
DLLVM_TARGETS_TO_BUILD="Native"
cmake --build . --target check-mlir # (take cap of tea, wait for 1h)
# done!

# on Windows (before building!) need to run regedit, go to path:
# HKEY_LOCAL_MACHINE\SYSTEM\CurrentControlSet\Control\FileSystem\
# set LongPathsEnabled = 1 and reboot the OS!
# Instead -G Ninja we use something like -G "Visual Studio 17 2022"
```


What is the hardest to me?

Middle end

Type inference – undecidable for many type systems
(in dataset ManyType4Py only 60-70% types are inferred)

Dialect design is an art (and, of course, TableGen and C++)

